

Study Guide

Expect a **subset** of questions drawn from the following topics.

Linear Regression Concepts

- A conceptual question describing and interpreting the effect of mixing continuous and categorical variables in linear regression, as discussed in Chapter 3 of ISLP.
- A question requiring you to list **at least five assumptions** associated with linear regression.
- A question on **heteroscedasticity**: what it is, why it is problematic in linear regression, and how it affects results and inference.
- A question in which you are given a **nonlinear relationship** and asked to linearize it by defining engineered features and writing the resulting linear regression model in the transformed space (e.g. polynomial regression).

Logistic Regression Theory

appendix7

1. Given the equation for **categorical cross-entropy**, show that it reduces to **binary cross-entropy** when ($C = 2$) (number of classes).
$$-\sum_{i=1}^n [y_i \log(p_i) + (1 - y_i) \log(1 - p_i)] \rightarrow \text{cross-entropy function}$$

appendix1

2. Work through a **loss function calculation** (either BCE or CCE) for a small toy classification dataset.

appendix2

3. Starting from the logit formulation, derive an explicit expression for ($p(X)$):

$$\log \left(\frac{p(X)}{1 - p(X)} \right) = \beta_0 + \beta^T X \quad (1)$$

appendix3

4. Compute the derivative of the **sigmoid function**.
5. Label a sketch of the sigmoid function given different values of fitting parameters (e.g., match parameter values to the appropriate curve).
6. **Bonus:** Derive an analytical expression for the derivative of the BCE loss with respect to the parameters of a univariate logistic regression model.

Interpreting Logistic Regression Results

Be familiar with the following aspects of interpreting logistic regression output. Expect **conceptual multiple-answer** or **short-answer** questions.

1. Estimated Coefficients ($\hat{\beta}$)

appendix4

- Represent how each predictor influences the outcome on the **log-odds** scale
- Define the fitted model and encode both the direction (increase/decrease) and strength of each effect
- A one-unit increase in (X_j) changes the log-odds of ($Y = 1$) by ($\hat{\beta}_j$), holding other variables fixed

2. Standard Errors

appendix5

- Quantify the **uncertainty** in coefficient estimates (derived from the observed Fisher information)
- Form the basis for confidence intervals and hypothesis tests
- Large standard errors suggest unstable estimates due to multicollinearity, small sample size, or separation

Interpretation: If we repeated the study many times, 95% of the time the true coefficient would fall in this interval.



3. Wald Z-Statistics and P-values

- Test the null hypothesis: $H_0 : \beta_j = 0$
- Measure how many standard errors the estimated coefficient is from zero
- Small p-values indicate statistically significant evidence of a nonzero effect

4. Odds Ratios and Confidence Intervals



- Obtained by exponentiating coefficients: $e^{\hat{\beta}_j}$
- Provide a more intuitive, multiplicative interpretation of effects on the odds scale
- An odds ratio (> 1) increases the odds of ($Y = 1$); an odds ratio (< 1) decreases the odds

appendix6

5. Confidence Intervals for Predicted Probabilities

- Quantify uncertainty in the predicted probability $\hat{p}(X)$ for specific predictor values
- Distinguish point predictions from uncertainty-aware probability estimates
- Narrow intervals indicate stable predictions; wide intervals often occur near separation or extrapolation

Example 1: Binary Cross-Entropy (BCE)

Setup

- Binary classification: $y \in \{0, 1\}$
- Loss for one observation:

$$\text{BCE}(y, p) = -[y \log(p) + (1 - y) \log(1 - p)]$$

Toy dataset (3 observations)

Observation	True label y	Predicted prob $p = P(y = 1)$
1	1	0.9
2	0	0.3
3	1	0.6

Step-by-step loss calculation

Observation 1

$$\text{Loss}_1 = -\log(0.9) = 0.105$$

Observation 2

$$\text{Loss}_2 = -\log(1 - 0.3) = -\log(0.7) = 0.357$$

Observation 3

$$\text{Loss}_3 = -\log(0.6) = 0.511$$

Total and average loss

Total BCE = $0.105 + 0.357 + 0.511 = 0.973$

Average BCE = $\frac{0.973}{3} = 0.324$

Example 2: Categorical Cross-Entropy (CCE)

Setup

- 3 classes: $C = 3$
- One-hot encoded labels
- Loss for one observation:

$$\text{CCE} = -\sum_{c=1}^C y_c \log(p_c)$$

Toy dataset (2 observations)

Observation 1

- True class: **Class 2**
- One-hot label:

$$y = [0, 1, 0]$$

- Predicted probabilities:

$$p = [0.2, 0.5, 0.3]$$

Loss:

$$\text{Loss}_1 = -\log(0.5) = 0.693$$

Observation 2

- True class: **Class 1**
- One-hot label:

$$y = [1, 0, 0]$$

- Predicted probabilities:

$$p = [0.7, 0.2, 0.1]$$

Loss:

$$\text{Loss}_2 = -\log(0.7) = 0.357$$

Total and average loss

Total CCE = $0.693 + 0.357 = 1.050$

appendix2

Given (logit formulation)

$$\log\left(\frac{p(x)}{1 - p(x)}\right) = \beta_0 + \beta^T x$$

where:

- $p(x) = P(Y = 1 \mid X = x)$
- $\beta^T x$ is the dot product of coefficients and predictors

Step 1: Exponentiate both sides

$$\frac{p(x)}{1 - p(x)} = e^{\beta_0 + \beta^T x}$$

Step 2: Solve for $p(x)$

Multiply both sides by $1 - p(x)$:

$$p(x) = (1 - p(x))e^{\beta_0 + \beta^T x}$$

Expand:

$$p(x) = e^{\beta_0 + \beta^T x} - p(x)e^{\beta_0 + \beta^T x}$$

Step 3: Collect $p(x)$ terms

$$p(x)(1 + e^{\beta_0 + \beta^T x}) = e^{\beta_0 + \beta^T x}$$

Step 4: Final expression

$$p(x) = \frac{e^{\beta_0 + \beta^T x}}{1 + e^{\beta_0 + \beta^T x}}$$

Equivalent (often preferred) form

$$p(x) = \frac{1}{1 + e^{-(\beta_0 + \beta^T x)}}$$

appendix3

$$\frac{d}{dz} \sigma(z)$$

2 Step 1: Rewrite for easier differentiation

$$\sigma(z) = (1 + e^{-z})^{-1}$$

3 Step 2: Use the chain rule

$$\frac{d}{dz} (1 + e^{-z})^{-1} = -(1 + e^{-z})^{-2} \cdot \frac{d}{dz} (1 + e^{-z})$$

Derivative of $1 + e^{-z}$ is:

$$\frac{d}{dz} (1 + e^{-z}) = -e^{-z}$$

So we get:

$$\frac{d}{dz} \sigma(z) = -(1 + e^{-z})^{-2} \cdot (-e^{-z}) = \frac{e^{-z}}{(1 + e^{-z})^2}$$

4 Step 3: Express in terms of $\sigma(z)$

Notice:

$$\sigma(z) = \frac{1}{1 + e^{-z}} \implies 1 - \sigma(z) = \frac{e^{-z}}{1 + e^{-z}}$$

Then:

$$\frac{e^{-z}}{(1 + e^{-z})^2} = \sigma(z) \cdot (1 - \sigma(z))$$

Final derivative

$$\frac{d}{dz} \sigma(z) = \sigma(z) (1 - \sigma(z))$$

Example 1: Continuous predictor (StudyHours)

Model:

$$\log \left(\frac{p(Y = 1)}{1 - p(Y = 1)} \right) = -2.0 + 0.8 \cdot \text{StudyHours}$$

Interpretation:

- A 1-unit increase in StudyHours increases the log-odds of $Y = 1$ by 0.8, holding other variables fixed.
- Exponentiating:

$$e^{0.8} \approx 2.23$$

- This means a one-hour increase in StudyHours multiplies the odds of $Y = 1$ by 2.23.

Example 2: Binary predictor (Smoker)

Model:

$$\log \left(\frac{p(Y = 1)}{1 - p(Y = 1)} \right) = -1.2 + 1.5 \cdot \text{Smoker}$$

(Smoker = 1 if smoker, 0 if non-smoker)

Interpretation:

- Being a smoker (vs non-smoker) increases the log-odds of $Y = 1$ by 1.5, holding other variables fixed.
- Exponentiating:

$$e^{1.5} \approx 4.48$$

- This means smokers have 4.48 times higher odds of $Y = 1$ compared to non-smokers.

 Quiz Question (Standard Errors in Logistic Regression)

You fit a logistic regression model and obtain the following output for a predictor X :

Coefficient	Std. Error	z-stat	p-value
0.90	0.15	6.00	< 0.001

Question:

What does the standard error of 0.15 tell you about the coefficient estimate, and how should it be interpreted?

 Correct Answer (Conceptual)

The standard error of 0.15 quantifies the uncertainty in the coefficient estimate of 0.90. It measures how much the estimated coefficient would vary across repeated samples. A relatively small standard error indicates that the coefficient estimate is precise and stable. Since the coefficient is much larger than its standard error, the estimate is statistically significant.

3 Example 1: Continuous predictor

Suppose the model is:

$$\log\left(\frac{p}{1-p}\right) = -2 + 0.8 \cdot \text{StudyHours}$$

- **Predict** for StudyHours = 3:

$$\hat{z} = -2 + 0.8 \cdot 3 = 0.4$$

- **Logistic transformation:**

$$\hat{p} = \frac{1}{1 + e^{-0.4}} \approx 0.60$$

- Suppose **SE** of $\hat{z} = 0.3$, then 95% CI for $\hat{z}$:

$$0.4 \pm 1.96 \cdot 0.3 = [-0.188, 0.988]$$

- **Transform back to probability:**

$$p_{\text{CI}} = \left[\frac{1}{1 + e^{0.188}}, \frac{1}{1 + e^{-0.988}} \right] \approx [0.46, 0.73]$$

Interpretation:

For 3 study hours, the predicted probability of passing is 0.60, but accounting for uncertainty, the true probability is likely between 0.46 and 0.73.

4 Example 2: Binary predictor

Model:

$$\log\left(\frac{p}{1-p}\right) = -1.0 + 1.2 \cdot \text{Smoker}$$

- **Point prediction:** Smoker = 1:

$$\hat{z} = -1.0 + 1.2 \cdot 1 = 0.2 \Rightarrow \hat{p} \approx 0.55$$

- SE of $\hat{z} = 0.4 \rightarrow$ 95% CI for log-odds:

$$0.2 \pm 1.96 \cdot 0.4 = [-0.584, 0.984]$$

- **Transform back to probability:**

$$p_{\text{CI}} = [0.36, 0.73]$$

1 Start with categorical cross-entropy (general form)

For C classes:

$$\mathcal{L}_{\text{cat}} = - \sum_{c=1}^C y_c \log(\hat{p}_c)$$

where:

- y_c is a one-hot encoded label
- $\hat{p}_c = P(Y = c)$

2 Set the number of classes to $C = 2$

$$\mathcal{L} = - [y_1 \log(\hat{p}_1) + y_2 \log(\hat{p}_2)]$$

5 Final result: binary cross-entropy

$$\mathcal{L}_{\text{binary}} = - [y \log(\hat{p}) + (1 - y) \log(1 - \hat{p})]$$

This is exactly the **binary cross-entropy** loss used in logistic regression.