

Integrated Gradients (IG)

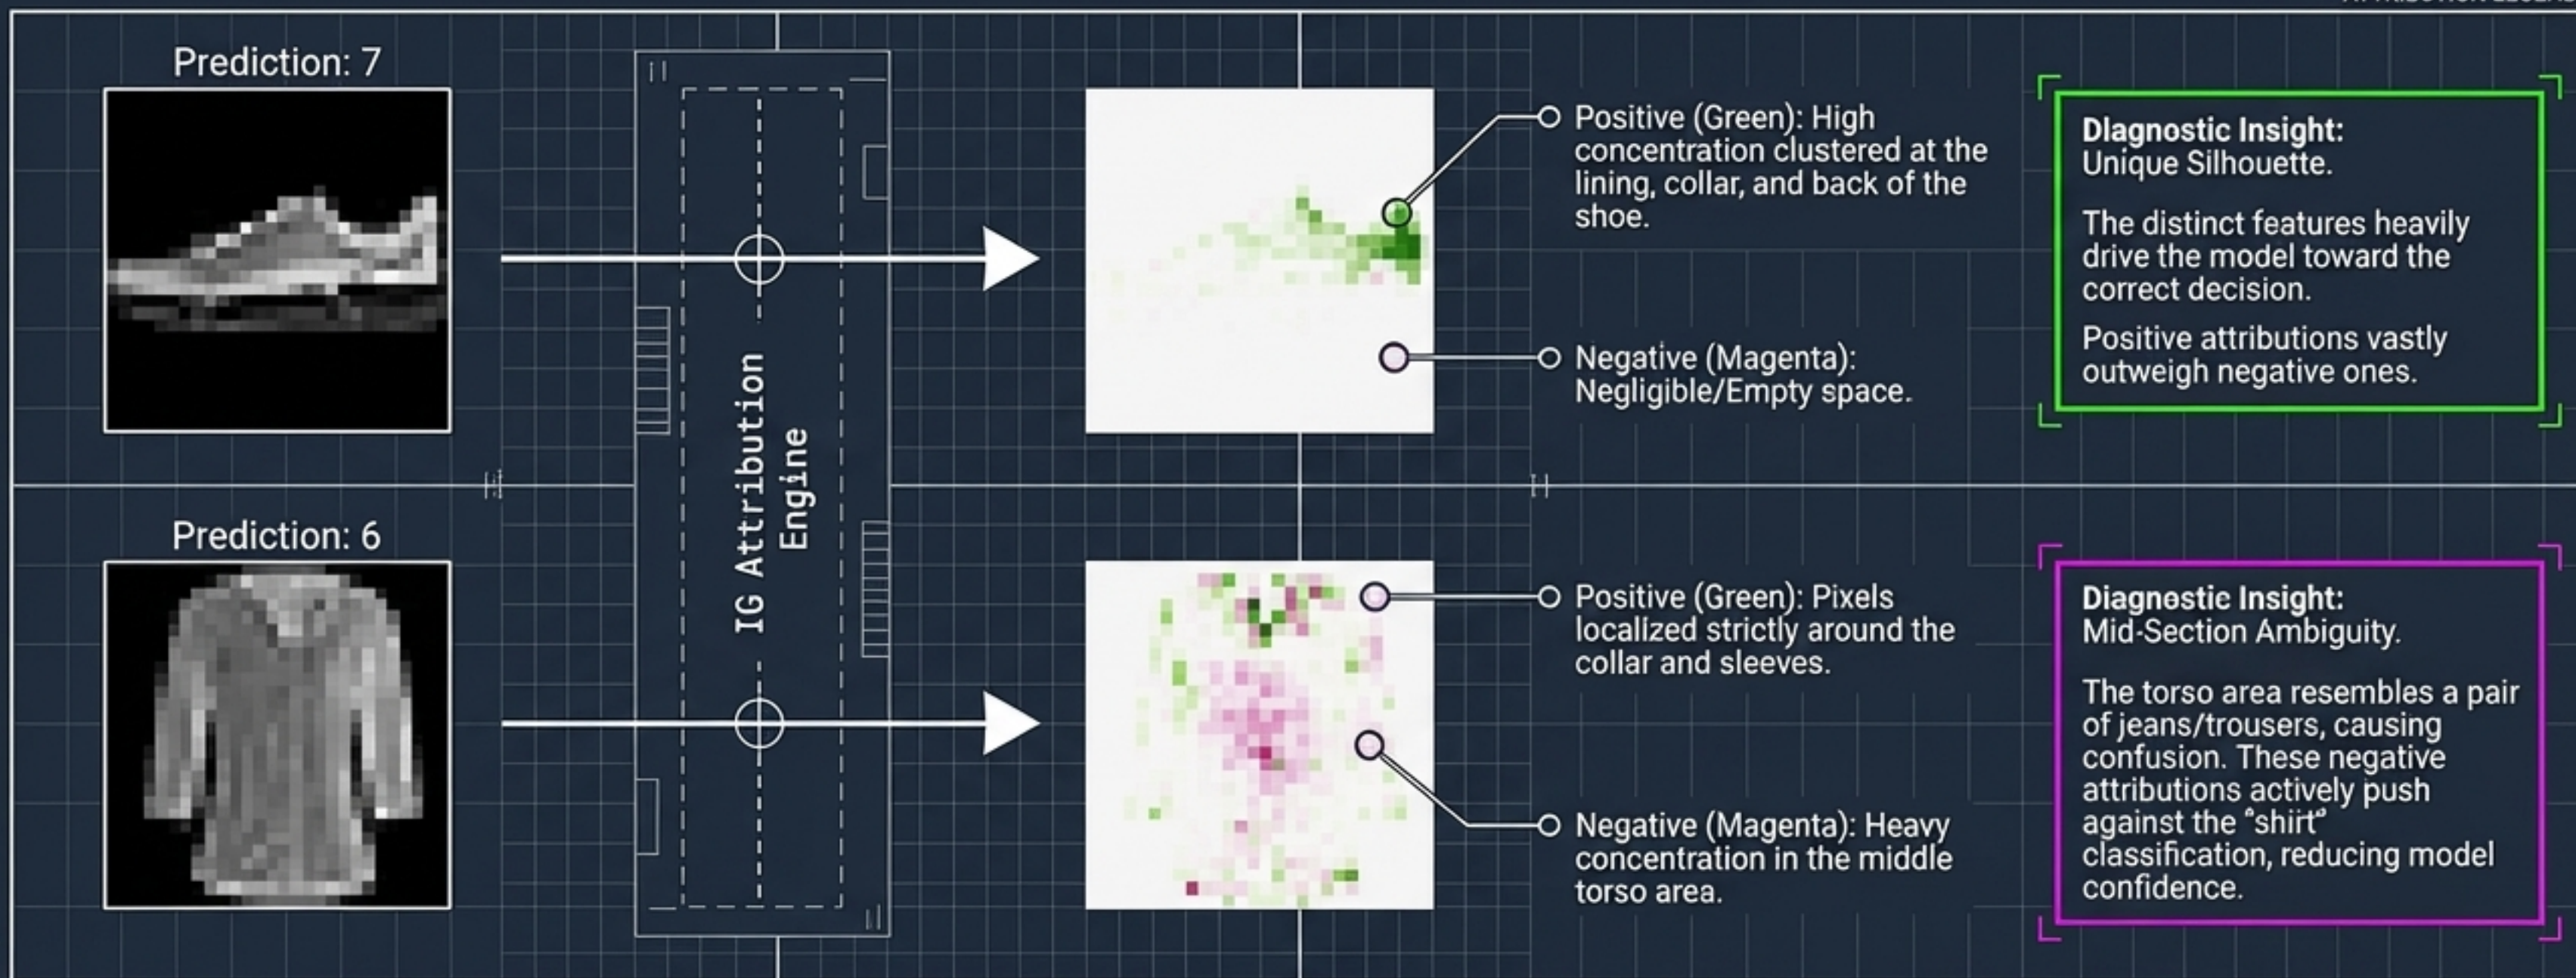
Local XAI Feature Attribution Schematic

A local Explainable AI (XAI) technique that assigns an importance value to each input feature using the gradients of the model's output. It opens the "black box" by revealing exactly which specific features influenced an individual prediction, highlighting opposing forces.

Positive Attributions: Features driving toward the decision.

Negative Attributions: Features driving against the decision.

ATTRIBUTION LEGEND



Captum

Official interpretability toolkit utilizing path-integrated gradients to understand important neurons and layers within PyTorch models.



Tf-explain

Integration for Keras API (TF v2.0 workflow) generating visual explanations and heatmaps for model predictions.